

at-5114-ai-in-hc-lp-diab-ch

March 6, 2024

[13]: IMPORTING LIBRARIES:

```
[53]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.plotting import scatter_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.feature_selection import RFE
from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score
import warnings

warnings.filterwarnings('ignore')
%matplotlib inline
```

[]: DATASET UPLOAD:

```
[54]: diabetes_df = pd.read_csv('C:\\Users\\HP\\Downloads\\archive (1)\\diabetes.csv')
diabetes_df.head()
```

```
[54]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0

4 2.288 33 1

```
[ ]: EXPLORATORY DATA ANALYSIS:
```

```
[55]: diabetes_df.columns
```

```
[55]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
         'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
        dtype='object')
```

```
[ ]: INFORMATION ABOUT DATASET:
```

```
[56]: diabetes_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 768 entries, 0 to 767  
Data columns (total 9 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   Pregnancies            768 non-null   int64  
1   Glucose                768 non-null   int64  
2   BloodPressure          768 non-null   int64  
3   SkinThickness          768 non-null   int64  
4   Insulin                768 non-null   int64  
5   BMI                    768 non-null   float64  
6   DiabetesPedigreeFunction 768 non-null   float64  
7   Age                    768 non-null   int64  
8   Outcome                768 non-null   int64  
dtypes: float64(2), int64(7)  
memory usage: 54.1 KB
```

```
[ ]: CHECKING THE NULL VALUES IN DATASET:
```

```
[57]: diabetes_df.isnull().head(10)
```

```
[57]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \  
0         False   False         False         False   False  False  
1         False   False         False         False   False  False  
2         False   False         False         False   False  False  
3         False   False         False         False   False  False  
4         False   False         False         False   False  False  
5         False   False         False         False   False  False  
6         False   False         False         False   False  False  
7         False   False         False         False   False  False  
8         False   False         False         False   False  False  
9         False   False         False         False   False  False
```

	DiabetesPedigreeFunction	Age	Outcome
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
5	False	False	False
6	False	False	False
7	False	False	False
8	False	False	False
9	False	False	False

```
[17]: diabetes_df.isnull().sum()
```

```
[17]: Pregnancies          0
      Glucose              0
      BloodPressure        0
      SkinThickness        0
      Insulin              0
      BMI                  0
      DiabetesPedigreeFunction  0
      Age                  0
      Outcome              0
      dtype: int64
```

```
[ ]: "In the given code snippet, the approach to handling missing values diverges
      ↪from traditional null-checking methods. Rather than explicitly checking for
      ↪null values using isnull() and calculating their sum, the code immediately
      ↪addresses potential missing data by creating a copy of the original dataset
      ↪(diabetes_df_copy). This copy is then subjected to imputation techniques
      ↪without initially distinguishing between explicit nulls and 0 values. The
      ↪specified columns - 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      ↪and 'BMI' - undergo mean imputation directly through fillna(diabetes_df_copy.
      ↪mean()). Consequently, missing or zero values are replaced with the mean of
      ↪each respective column. The subsequent display of NaN counts through
      ↪print(diabetes_df_copy.isnull().sum()) provides insights into the
      ↪effectiveness of the imputation process in handling missing or zero values
      ↪within the dataset."
```

```
[58]: import numpy as np
      import pandas as pd

      # Load the dataset
      diabetes_df = pd.read_csv('C:\\Users\\HP\\Downloads\\archive (1)\\diabetes.csv')

      # Define columns to replace and create a copy of the dataframe
```

```

columns_to_replace = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']
diabetes_df_copy = diabetes_df.copy(deep=True)

# Alternatively, consider using imputation techniques for missing values
# Using mean imputation for the specified columns
diabetes_df_copy.fillna(diabetes_df_copy.mean(), inplace=True)

# Displaying the Count of NaNs after imputation
print(diabetes_df_copy.isnull().sum())

```

```


Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age              0
Outcome           0
dtype: int64

```

```
[ ]: DATA VISUALIZATION:
```

```
[ ]: DATA DISTRIBUTION PLOTS BEFORE REMOVING THE NULL VALUES :
```

```
[59]: p = diabetes_df.hist(figsize = (20,20))
```




[]: Now we will be imputing the mean value of the column to each missing value of `Glucose` that particular column.

```
[27]: diabetes_df_copy['Glucose'].fillna(diabetes_df_copy['Glucose'].mean(), inplace=True)
diabetes_df_copy['BloodPressure'].fillna(diabetes_df_copy['BloodPressure'].mean(), inplace=True)
diabetes_df_copy['SkinThickness'].fillna(diabetes_df_copy['SkinThickness'].median(), inplace=True)
diabetes_df_copy['Insulin'].fillna(diabetes_df_copy['Insulin'].median(), inplace=True)
diabetes_df_copy['BMI'].fillna(diabetes_df_copy['BMI'].median(), inplace=True)
```


```
[ ]: DATA DISTRIBUTION PLOTS AFTER REMOVING THE NULL VALUES:
```

```
[60]: # Data Distribution Plots after imputation  
p = diabetes_df_copy.hist(figsize=(20, 20))
```



```
[ ]: FEATURES CORRELATION:
```

```
[61]: plt.figure(figsize=(12, 10))  
correlation_matrix = diabetes_df.corr()  
sns.heatmap(correlation_matrix, annot=True, cmap='RdYlGn')  
plt.show()
```



```
[ ]: DATA SCALING:
```

```
[ ]: Let's check the data before scaling it
```

```
[62]: diabetes_df_copy.head()
```

```
[62]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6      148           72           35         0  33.6
1           1       85           66           29         0  26.6
2           8      183           64            0         0  23.3
3           1       89           66           23        94  28.1
4           0      137           40           35       168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627      50         1
```

1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

[]: AFTER STANDARD SCALING:

```
[63]: columns_to_scale = ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
    ↪ 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']

sc_X = StandardScaler()
X = pd.DataFrame(sc_X.fit_transform(diabetes_df_copy[columns_to_scale]),
    ↪ columns=columns_to_scale)
X.head()
```

```
[63]:   Pregnancies   Glucose  BloodPressure  SkinThickness   Insulin   BMI \
0      0.639947  0.848324      0.149641      0.907270 -0.692891  0.204013
1     -0.844885 -1.123396     -0.160546      0.530902 -0.692891 -0.684422
2      1.233880  1.943724     -0.263941     -1.288212 -0.692891 -1.103255
3     -0.844885 -0.998208     -0.160546      0.154533  0.123302 -0.494043
4     -1.141852  0.504055     -1.504687      0.907270  0.765836  1.409746

      DiabetesPedigreeFunction   Age
0              0.468492  1.425995
1             -0.365061 -0.190672
2              0.604397 -0.105584
3             -0.920763 -1.041549
4              5.484909 -0.020496
```

[]: MODEL BUILDING USING ML ALGORITHMS:

[]: RANDOM FOREST

[]: DATA SPLITTING: (FEATURE SELECTION)

```
[64]: # Feature selection using Random Forest
# Specify the features and target variable
X = diabetes_df_copy.drop('Outcome', axis=1)
y = diabetes_df_copy['Outcome']

# Perform the train-test split
random_state_value = 7 # You can change this value if needed
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
    ↪ random_state=random_state_value)

# Create and train the Random Forest Classifier for feature selection
rfc_feature_selection = RandomForestClassifier(n_estimators=200)
```



```

rfc_feature_selection.fit(X_train, y_train)

# Extract feature importances
feature_importances = rfc_feature_selection.feature_importances_

# Create a DataFrame to display feature importances
feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importances
})


# Sort features by importance in descending order
feature_importance_df = feature_importance_df.sort_values(by='Importance',
    ↪ascending=False)

# Plotting the feature importances
plt.figure(figsize=(10, 8))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df,
    ↪palette='viridis')
plt.title('Feature Importances')
plt.show()

# Select top features based on importance
threshold = 0.02 # Adjust the threshold as needed
selected_features = feature_importance_df.
    ↪loc[feature_importance_df['Importance'] > threshold]['Feature'].values

# Use only the selected features for training and testing
X_train_selected = X_train[selected_features]
X_test_selected = X_test[selected_features]

```



```
[65]: # MODEL BUILDING USING ML ALGORITHMS:
# RANDOM FOREST:
# Create and train the RandomForestClassifier with selected features
rfc = RandomForestClassifier(n_estimators=200)
rfc.fit(X_train_selected, y_train)

# Predict on the test set
predictions_rfc = rfc.predict(X_test_selected)

# Calculate and print the accuracy score
accuracy_rfc = accuracy_score(y_test, predictions_rfc)
print("Random Forest Accuracy Score =", format(accuracy_rfc))
```

Random Forest Accuracy Score = 0.7716535433070866

```
[ ]: CLASSIFICATION REPORT & CONFUSION MATRIX:
```

```
[66]: # Confusion Matrix:
cm_rfc = confusion_matrix(y_test, predictions_rfc)


# Create a heatmap for the confusion matrix
plt.figure(figsize=(8, 6))
```

```

sns.heatmap(cm_rfc, annot=True, fmt='d', cmap='Blues', linewidths=.5,
            cbar=False,
            xticklabels=['Predicted 0', 'Predicted 1'], yticklabels=['Actual_
            ↪0', 'Actual 1'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for Random Forest')
plt.show()

# Classification Report:
classification_report_rfc = classification_report(y_test, predictions_rfc)
print("Classification Report for Random Forest:\n", classification_report_rfc)

```



Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.81	0.85	0.83	162
1	0.70	0.64	0.67	92

accuracy			0.77	254
macro avg	0.75	0.74	0.75	254
weighted avg	0.77	0.77	0.77	254

[]: DECISION TREE:

```
[67]: # Create and train the DecisionTreeClassifier for feature selection
dtree_feature_selection = DecisionTreeClassifier()
dtree_feature_selection.fit(X_train, y_train)

# Extract feature importances
feature_importances_dtree = dtree_feature_selection.feature_importances_

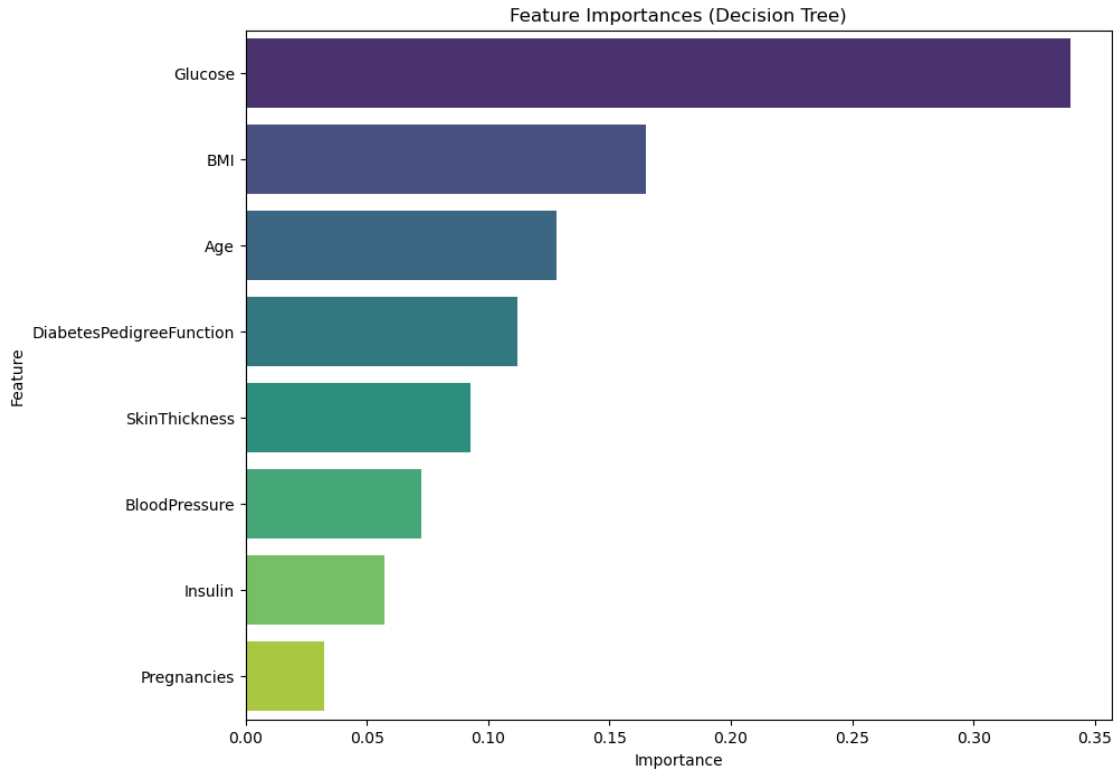
# Create a DataFrame to display feature importances
feature_importance_df_dtree = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importances_dtree
})

# Sort features by importance in descending order
feature_importance_df_dtree = feature_importance_df_dtree.
    ↪sort_values(by='Importance', ascending=False)

# Plotting the feature importances for Decision Tree
plt.figure(figsize=(10, 8))
sns.barplot(x='Importance', y='Feature', data=feature_importance_df_dtree,
    ↪palette='viridis')
plt.title('Feature Importances (Decision Tree)')
plt.show()

# Select top features based on importance
threshold_dtree = 0.02 # Adjust the threshold as needed
selected_features_dtree = feature_importance_df_dtree.loc[
    feature_importance_df_dtree['Importance'] > threshold_dtree]['Feature'].
    ↪values

# Use only the selected features for training and testing for Decision Tree
X_train_selected_dtree = X_train[selected_features_dtree]
X_test_selected_dtree = X_test[selected_features_dtree]
```



```
[68]: from sklearn.tree import DecisionTreeClassifier
```

```
# Create and train the DecisionTreeClassifier
dtree = DecisionTreeClassifier()
dtree.fit(X_train, y_train)
```

```
[68]: DecisionTreeClassifier()
```

```
[ ]: CHECKING THE ACCURACY SCORE ON TESTING SET:
```

```
[69]: # Create and train the DecisionTreeClassifier with selected features
```

```
dtree = DecisionTreeClassifier()
dtree.fit(X_train_selected_dtree, y_train)

# Predict on the test set for Decision Tree
predictions_dtree = dtree.predict(X_test_selected_dtree)
```

```
# Calculate and print the accuracy score for Decision Tree
accuracy_dtree = accuracy_score(y_test, predictions_dtree)
print("Decision Tree Accuracy Score =", format(accuracy_dtree))
```

Decision Tree Accuracy Score = 0.7165354330708661

```
[ ]: CLASSIFICATION REPORT & CONFUSION MATRIX:
```

```
[70]: # Assuming 'X_test' contains the testing features with the same columns as
      ↪ 'X_train'


      # Train Decision Tree Classifier
      dtree = DecisionTreeClassifier()
      dtree.fit(X_train, y_train)

      # Predict on the test set with the same features used for training
      predictions_dtree = dtree.predict(X_test)

      # Confusion Matrix:
      cm_dtree = confusion_matrix(y_test, predictions_dtree)

      # Create a heatmap for the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cm_dtree, annot=True, fmt='d', cmap='Blues', linewidths=.5,
      ↪ cbar=False,
      ↪ xticklabels=['Predicted 0', 'Predicted 1'], yticklabels=['Actual_
      ↪ 0', 'Actual 1'])
      plt.xlabel('Predicted Label')
      plt.ylabel('True Label')
      plt.title('Confusion Matrix for Decision Tree')
      plt.show()

      # Classification Report:
      classification_report_dtree = classification_report(y_test, predictions_dtree)
      print("Classification Report for Decision Tree:\n", classification_report_dtree)
```



Classification Report for Decision Tree:

	precision	recall	f1-score	support
0	0.77	0.78	0.78	162
1	0.61	0.60	0.60	92
accuracy			0.72	254
macro avg	0.69	0.69	0.69	254
weighted avg	0.72	0.72	0.72	254

[]: SUPPORT VECTOR MACHINE (SVM):

```
[71]: # Feature selection using Recursive Feature Elimination (RFE) for Support
      ↪ Vector Machines (SVM)
from sklearn.svm import SVC, LinearSVC
from sklearn.feature_selection import RFE
from sklearn.model_selection import train_test_split
```

```

# Specify the features and target variable
X = diabetes_df_copy.drop('Outcome', axis=1)
y = diabetes_df_copy['Outcome']

# Perform the train-test split
random_state_value = 7 # You can change this value if needed
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
    ↪random_state=random_state_value)

# Create and train the LinearSVC model
linear_svc_model = LinearSVC()
linear_svc_model.fit(X_train, y_train)

# Use Recursive Feature Elimination (RFE) for feature selection
num_features_to_select = 5 # Adjust the number of features as needed
rfe = RFE(linear_svc_model, n_features_to_select=num_features_to_select)
rfe.fit(X_train, y_train)

# Extract selected features
selected_features = X_train.columns[rfe.support_]

# Use only the selected features for training and testing
X_train_selected = X_train[selected_features]
X_test_selected = X_test[selected_features]

# Display the selected features
print("Selected Features:", selected_features)

```

```

Selected Features: Index(['Pregnancies', 'BloodPressure', 'SkinThickness',
'BMI',
'DiabetesPedigreeFunction'],
dtype='object')

```

```

[72]: from sklearn.svm import SVC
      from sklearn import metrics

      # Assuming 'X_test' contains the testing features
      svc_model = SVC()
      svc_model.fit(X_train, y_train)

```

```
[72]: SVC()
```

```

[76]: from sklearn.metrics import accuracy_score

      # Assuming 'y_test' contains the true labels and 'svc_pred' contains the
      ↪predicted labels
      svc_pred = svc_model.predict(X_test)

```



```

# Calculate Accuracy Score for SVC
accuracy_score_svc = accuracy_score(y_test, svc_pred)

# Display the Accuracy Score
print("Accuracy Score for SVC:", accuracy_score_svc)

```

Accuracy Score for SVC: 0.7480314960629921

```

[78]: from sklearn.metrics import confusion_matrix, classification_report

# Assuming 'y_test' contains the true labels and 'svc_pred' contains the
# predicted labels
svc_pred = svc_model.predict(X_test)

# Confusion Matrix for SVC
cm_svc = confusion_matrix(y_test, svc_pred)
print("Confusion Matrix for SVC:\n", cm_svc)

# Create a heatmap for the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cm_svc, annot=True, fmt='d', cmap='Blues', linewidths=.5,
            cbar=False,
            xticklabels=['Predicted 0', 'Predicted 1'], yticklabels=['Actual
            ↪0', 'Actual 1'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for SVC')
plt.show()

# Classification Report for SVC
classification_report_svc = classification_report(y_test, svc_pred)
print("Classification Report for SVC:\n", classification_report_svc)


```

Confusion Matrix for SVC:

```

[[145  17]
 [ 47  45]]

```



Classification Report for SVC:

	precision	recall	f1-score	support
0	0.76	0.90	0.82	162
1	0.73	0.49	0.58	92
accuracy			0.75	254
macro avg	0.74	0.69	0.70	254
weighted avg	0.74	0.75	0.73	254

```
[84]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Example: Replace these with your actual metric values
model_names = ['Random Forest', 'Decision Tree', 'LinearSVC']
accuracy_list = [0.77, 0.72, 0.75]
precision_list = [0.81, 0.77, 0.76]
```

```

recall_list = [0.85, 0.78, 0.90]
f1_score_list = [0.83, 0.78, 0.82]

# Create a DataFrame for model comparison metrics
model_comparison_metrics = pd.DataFrame({
    'Model': model_names,
    'Accuracy': accuracy_list,
    'Precision': precision_list,
    'Recall': recall_list,
    'F1 Score': f1_score_list
})

# Create a heatmap for model comparison metrics
plt.figure(figsize=(10, 8))
sns.heatmap(model_comparison_metrics.set_index('Model')[['Accuracy', 'Precision', 'Recall', 'F1 Score']], annot=True, cmap='Blues', linewidths=.5, cbar=False)
plt.xlabel('Metrics')
plt.ylabel('Model')
plt.title('Model Comparison Metrics Heatmap')
plt.show()

```

